



# Memory Management

## Overview

Memory management in Python involves a private heap containing all Python objects and data structures. The management of this private heap is ensured internally by the *Python memory manager*. The Python memory manager has different components which deal with various dynamic storage management aspects, like sharing, segmentation, preallocation or caching.

At the lowest level, a raw memory allocator ensures that there is enough room in the private heap for storing all Python-related data by interacting with the memory manager of the operating system. On top of the raw memory allocator, several object-specific allocators operate on the same heap and implement distinct memory management policies adapted to the peculiarities of every object type. For example, integer objects are managed differently within the heap than strings, tuples or dictionaries because integers imply different storage requirements and speed/space tradeoffs. The Python memory manager thus delegates some of the work to the object-specific allocators, but ensures that the latter operate within the bounds of the private heap.

It is important to understand that the management of the Python heap is performed by the interpreter itself and that the user has no control over it, even if they regularly manipulate object pointers to memory blocks inside that heap. The allocation of heap space for Python objects and other internal buffers is performed on demand by the Python memory manager through the Python/C API functions listed in this document.

To avoid memory corruption, extension writers should never try to operate on Python objects with the functions exported by the C library: `malloc()`, `calloc()`, `realloc()` and `free()`. This will result in mixed calls between the C allocator and the Python memory manager with fatal consequences, because they implement different algorithms and operate on different heaps. However, one may safely allocate and release memory blocks with the C library allocator for individual purposes, as shown in the following example:

```
PyObject *res;
char *buf = (char *) malloc(BUFSIZ); /* for I/O */

if (buf == NULL)
    return PyErr_NoMemory();
...Do some I/O operation involving buf...
res = PyBytes_FromString(buf);
free(buf); /* malloc'ed */
return res;
```

In this example, the memory request for the I/O buffer is handled by the C library allocator. The Python memory manager is involved only in the allocation of the bytes object returned as a result.

In most situations, however, it is recommended to allocate memory from the Python heap specifically because the latter is under control of the Python memory manager. For example, this is required when the interpreter is extended with new object types written in C. Another reason for using the Python heap is the desire to *inform* the Python memory manager about the memory needs of the extension module. Even when the requested memory is used exclusively for internal, highly-specific purposes, delegating all memory requests to the Python memory manager causes the interpreter to have a more accurate image of its memory footprint as a whole. Consequently, under certain circumstances, the Python memory manager may or may not trigger appropriate actions, like garbage collection, memory compaction or other preventive procedures. Note that by using the C



library allocator as shown in the previous example, the allocated memory for the I/O buffer escapes completely the Python memory manager.

**See also:** The `PYTHONMALLOC` environment variable can be used to configure the memory allocators used by Python.

The `PYTHONMALLOCSTATS` environment variable can be used to print statistics of the `pymalloc` memory allocator every time a new pymalloc object arena is created, and on shutdown.

## Allocator Domains

All allocating functions belong to one of three different “domains” (see also `PyMemAllocatorDomain`). These domains represent different allocation strategies and are optimized for different purposes. The specific details on how every domain allocates memory or what internal functions each domain calls is considered an implementation detail, but for debugging purposes a simplified table can be found at [here](#). There is no hard requirement to use the memory returned by the allocation functions belonging to a given domain for only the purposes hinted by that domain (although this is the recommended practice). For example, one could use the memory returned by `PyMem_RawMalloc()` for allocating Python objects or the memory returned by `PyObject_Malloc()` for allocating memory for buffers.

The three allocation domains are:

- Raw domain: intended for allocating memory for general-purpose memory buffers where the allocation *must* go to the system allocator or where the allocator can operate without the `GIL`. The memory is requested directly to the system.
- “Mem” domain: intended for allocating memory for Python buffers and general-purpose memory buffers where the allocation must be performed with the `GIL` held. The memory is taken from the Python private heap.
- Object domain: intended for allocating memory belonging to Python objects. The memory is taken from the Python private heap.

When freeing memory previously allocated by the allocating functions belonging to a given domain, the matching specific deallocating functions must be used. For example, `PyMem_Free()` must be used to free memory allocated using `PyMem_Malloc()`.

## Raw Memory Interface

The following function sets are wrappers to the system allocator. These functions are thread-safe, the `GIL` does not need to be held.

The `default raw memory allocator` uses the following functions: `malloc()`, `calloc()`, `realloc()` and `free()`; call `malloc(1)` (or `calloc(1, 1)`) when requesting zero bytes.

*New in version 3.4.*

`void *PyMem_RawMalloc(size_t n)`

Allocates *n* bytes and returns a pointer of type `void*` to the allocated memory, or `NULL` if the request fails.

Requesting zero bytes returns a distinct non-`NULL` pointer if possible, as if `PyMem_RawMalloc(1)` had been called instead. The memory will not have been initialized in any



void \***PyMem\_RawCalloc**(size\_t *nelem*, size\_t *elsize*)

Allocates *nelem* elements each whose size in bytes is *elsize* and returns a pointer of type void\* to the allocated memory, or NULL if the request fails. The memory is initialized to zeros.

Requesting zero elements or elements of size zero bytes returns a distinct non-NULL pointer if possible, as if `PyMem_RawCalloc(1, 1)` had been called instead.

*New in version 3.5.*

void \***PyMem\_RawRealloc**(void \**p*, size\_t *n*)

Resizes the memory block pointed to by *p* to *n* bytes. The contents will be unchanged to the minimum of the old and the new sizes.

If *p* is NULL, the call is equivalent to `PyMem_RawMalloc(n)`; else if *n* is equal to zero, the memory block is resized but is not freed, and the returned pointer is non-NULL.

Unless *p* is NULL, it must have been returned by a previous call to `PyMem_RawMalloc()`, `PyMem_RawRealloc()` or `PyMem_RawCalloc()`.

If the request fails, `PyMem_RawRealloc()` returns NULL and *p* remains a valid pointer to the previous memory area.

void **PyMem\_RawFree**(void \**p*)

Frees the memory block pointed to by *p*, which must have been returned by a previous call to `PyMem_RawMalloc()`, `PyMem_RawRealloc()` or `PyMem_RawCalloc()`. Otherwise, or if `PyMem_RawFree(p)` has been called before, undefined behavior occurs.

If *p* is NULL, no operation is performed.

## Memory Interface

The following function sets, modeled after the ANSI C standard, but specifying behavior when requesting zero bytes, are available for allocating and releasing memory from the Python heap.

The [default memory allocator](#) uses the [pymalloc memory allocator](#).

**Warning:** The [GIL](#) must be held when using these functions.

*Changed in version 3.6:* The default allocator is now `pymalloc` instead of `system malloc()`.

void \***PyMem\_Malloc**(size\_t *n*)

*Part of the Stable ABI.*

Allocates *n* bytes and returns a pointer of type void\* to the allocated memory, or NULL if the request fails.

Requesting zero bytes returns a distinct non-NULL pointer if possible, as if `PyMem_Malloc(1)` had been called instead. The memory will not have been initialized in any way.

void \***PyMem\_Calloc**(size\_t *nelem*, size\_t *elsize*)

*Part of the Stable ABI since version 3.7.*

Allocates *nelem* elements each whose size in bytes is *elsize* and returns a pointer of type void\*



Requesting zero elements or elements of size zero bytes returns a distinct non-NULL pointer if possible, as if `PyMem_Calloc(1, 1)` had been called instead.

*New in version 3.5.*

`void *PyMem_Realloc(void *p, size_t n)`

*Part of the Stable ABI.*

Resizes the memory block pointed to by *p* to *n* bytes. The contents will be unchanged to the minimum of the old and the new sizes.

If *p* is NULL, the call is equivalent to `PyMem_Malloc(n)`; else if *n* is equal to zero, the memory block is resized but is not freed, and the returned pointer is non-NULL.

Unless *p* is NULL, it must have been returned by a previous call to `PyMem_Malloc()`, `PyMem_Realloc()` or `PyMem_Calloc()`.

If the request fails, `PyMem_Realloc()` returns NULL and *p* remains a valid pointer to the previous memory area.

`void PyMem_Free(void *p)`

*Part of the Stable ABI.*

Frees the memory block pointed to by *p*, which must have been returned by a previous call to `PyMem_Malloc()`, `PyMem_Realloc()` or `PyMem_Calloc()`. Otherwise, or if `PyMem_Free(p)` has been called before, undefined behavior occurs.

If *p* is NULL, no operation is performed.

The following type-oriented macros are provided for convenience. Note that *TYPE* refers to any C type.

`TYPE *PyMem_New(TYPE, size_t n)`

Same as `PyMem_Malloc()`, but allocates `(n * sizeof(TYPE))` bytes of memory. Returns a pointer cast to `TYPE*`. The memory will not have been initialized in any way.

`TYPE *PyMem_Resize(void *p, TYPE, size_t n)`

Same as `PyMem_Realloc()`, but the memory block is resized to `(n * sizeof(TYPE))` bytes. Returns a pointer cast to `TYPE*`. On return, *p* will be a pointer to the new memory area, or NULL in the event of failure.

This is a C preprocessor macro; *p* is always reassigned. Save the original value of *p* to avoid losing memory when handling errors.

`void PyMem_Del(void *p)`

Same as `PyMem_Free()`.

In addition, the following macro sets are provided for calling the Python memory allocator directly, without involving the C API functions listed above. However, note that their use does not preserve binary compatibility across Python versions and is therefore deprecated in extension modules.

- `PyMem_MALLOC(size)`
- `PyMem_NEW(type, size)`
- `PyMem_REALLOC(ptr, size)`



- `PyMem_FREE(ptr)`
- `PyMem_DEL(ptr)`

## Object allocators

The following function sets, modeled after the ANSI C standard, but specifying behavior when requesting zero bytes, are available for allocating and releasing memory from the Python heap.

**Note:** There is no guarantee that the memory returned by these allocators can be successfully cast to a Python object when intercepting the allocating functions in this domain by the methods described in the [Customize Memory Allocators](#) section.

The [default object allocator](#) uses the [pymalloc memory allocator](#).

**Warning:** The [GIL](#) must be held when using these functions.

`void *PyObject_Malloc(size_t n)`

*Part of the Stable ABI.*

Allocates *n* bytes and returns a pointer of type `void*` to the allocated memory, or `NULL` if the request fails.

Requesting zero bytes returns a distinct non-`NULL` pointer if possible, as if `PyObject_Malloc(1)` had been called instead. The memory will not have been initialized in any way.

`void *PyObject_Calloc(size_t nelem, size_t elsize)`

*Part of the Stable ABI since version 3.7.*

Allocates *nelem* elements each whose size in bytes is *elsize* and returns a pointer of type `void*` to the allocated memory, or `NULL` if the request fails. The memory is initialized to zeros.

Requesting zero elements or elements of size zero bytes returns a distinct non-`NULL` pointer if possible, as if `PyObject_Calloc(1, 1)` had been called instead.

*New in version 3.5.*

`void *PyObject_Realloc(void *p, size_t n)`

*Part of the Stable ABI.*

Resizes the memory block pointed to by *p* to *n* bytes. The contents will be unchanged to the minimum of the old and the new sizes.

If *p* is `NULL`, the call is equivalent to `PyObject_Malloc(n)`; else if *n* is equal to zero, the memory block is resized but is not freed, and the returned pointer is non-`NULL`.

Unless *p* is `NULL`, it must have been returned by a previous call to [PyObject\\_Malloc\(\)](#), [PyObject\\_Realloc\(\)](#) or [PyObject\\_Calloc\(\)](#).

If the request fails, [PyObject\\_Realloc\(\)](#) returns `NULL` and *p* remains a valid pointer to the previous memory area.

`void PyObject_Free(void *p)`



Frees the memory block pointed to by *p*, which must have been returned by a previous call to `PyObject_Malloc()`, `PyObject_Realloc()` or `PyObject_Calloc()`. Otherwise, or if `PyObject_Free(p)` has been called before, undefined behavior occurs.

If *p* is `NULL`, no operation is performed.

## Default Memory Allocators

Default memory allocators:

Configuration	Name	PyMem_RawMalloc	PyMem_Malloc	PyObject_Malloc
Release build	"pymalloc"	malloc	pymalloc	pymalloc
Debug build	"pymalloc_debug"	malloc + debug	pymalloc + debug	pymalloc + debug
Release build, without pymalloc	"malloc"	malloc	malloc	malloc
Debug build, without pymalloc	"malloc_debug"	malloc + debug	malloc + debug	malloc + debug

Legend:

- Name: value for `PYTHONMALLOC` environment variable.
- malloc: system allocators from the standard C library, C functions: `malloc()`, `calloc()`, `realloc()` and `free()`.
- pymalloc: [pymalloc memory allocator](#).
- "+ debug": with [debug hooks on the Python memory allocators](#).
- "Debug build": [Python build in debug mode](#).

## Customize Memory Allocators

*New in version 3.4.*

type **PyMemAllocatorEx**

Structure used to describe a memory block allocator. The structure has the following fields:

Field	Meaning
<code>void *ctx</code>	user context passed as first argument
<code>void* malloc(void *ctx, size_t size)</code>	allocate a memory block
<code>void* calloc(void *ctx, size_t nelem, size_t elsize)</code>	allocate a memory block initialized with zeros
<code>void* realloc(void *ctx, void *ptr, size_t new_size)</code>	allocate or resize a memory block
<code>void free(void *ctx, void *ptr)</code>	free a memory block



*Changed in version 3.5:* The `PyMemAllocator` structure was renamed to `PyMemAllocatorEx` and a new `calloc` field was added.

### type **PyMemAllocatorDomain**

Enum used to identify an allocator domain. Domains:

#### **PYMEM\_DOMAIN\_RAW**

Functions:

- `PyMem_RawMalloc()`
- `PyMem_RawRealloc()`
- `PyMem_RawCalloc()`
- `PyMem_RawFree()`

#### **PYMEM\_DOMAIN\_MEM**

Functions:

- `PyMem_Malloc()`,
- `PyMem_Realloc()`
- `PyMem_Calloc()`
- `PyMem_Free()`

#### **PYMEM\_DOMAIN\_OBJ**

Functions:

- `PyObject_Malloc()`
- `PyObject_Realloc()`
- `PyObject_Calloc()`
- `PyObject_Free()`

void **PyMem\_GetAllocator**(`PyMemAllocatorDomain` *domain*, `PyMemAllocatorEx` \**allocator*)

Get the memory block allocator of the specified domain.

void **PyMem\_SetAllocator**(`PyMemAllocatorDomain` *domain*, `PyMemAllocatorEx` \**allocator*)

Set the memory block allocator of the specified domain.

The new allocator must return a distinct non-NULL pointer when requesting zero bytes.

For the `PYMEM_DOMAIN_RAW` domain, the allocator must be thread-safe: the [GIL](#) is not held when the allocator is called.

If the new allocator is not a hook (does not call the previous allocator), the `PyMem_SetupDebugHooks()` function must be called to reinstall the debug hooks on top on the new allocator.

void **PyMem\_SetupDebugHooks**(void)

Setup [debug hooks in the Python memory allocators](#) to detect memory errors.

## Debug hooks on the Python memory allocators

When Python is built in debug mode, the `PyMem_SetupDebugHooks()` function is called at the [Python preinitialization](#) to setup debug hooks on Python memory allocators to detect memory errors





The `PYTHONMALLOC` environment variable can be used to install debug hooks on a Python compiled in release mode (ex: `PYTHONMALLOC=debug`).

The `PyMem_SetupDebugHooks()` function can be used to set debug hooks after calling `PyMem_SetAllocator()`.

These debug hooks fill dynamically allocated memory blocks with special, recognizable bit patterns. Newly allocated memory is filled with the byte `0xCD` (`PYMEM_CLEANBYTE`), freed memory is filled with the byte `0xDD` (`PYMEM_DEADBYTE`). Memory blocks are surrounded by “forbidden bytes” filled with the byte `0xFD` (`PYMEM_FORBIDDENBYTE`). Strings of these bytes are unlikely to be valid addresses, floats, or ASCII strings.

Runtime checks:

- Detect API violations. For example, detect if `PyObject_Free()` is called on a memory block allocated by `PyMem_Malloc()`.
- Detect write before the start of the buffer (buffer underflow).
- Detect write after the end of the buffer (buffer overflow).
- Check that the `GIL` is held when allocator functions of `PYMEM_DOMAIN_OBJ` (ex: `PyObject_Malloc()`) and `PYMEM_DOMAIN_MEM` (ex: `PyMem_Malloc()`) domains are called.

On error, the debug hooks use the `tracemalloc` module to get the traceback where a memory block was allocated. The traceback is only displayed if `tracemalloc` is tracing Python memory allocations and the memory block was traced.

Let  $S = \text{sizeof}(\text{size\_t})$ .  $2 \cdot S$  bytes are added at each end of each block of  $N$  bytes requested. The memory layout is like so, where  $p$  represents the address returned by a `malloc`-like or `realloc`-like function ( $p[i:j]$  means the slice of bytes from  $*(p+i)$  inclusive up to  $*(p+j)$  exclusive; note that the treatment of negative indices differs from a Python slice):

$p[-2 \cdot S:-S]$

Number of bytes originally asked for. This is a `size_t`, big-endian (easier to read in a memory dump).

$p[-S]$

API identifier (ASCII character):

- 'r' for `PYMEM_DOMAIN_RAW`.
- 'm' for `PYMEM_DOMAIN_MEM`.
- 'o' for `PYMEM_DOMAIN_OBJ`.

$p[-S+1:0]$

Copies of `PYMEM_FORBIDDENBYTE`. Used to catch under- writes and reads.

$p[0:N]$

The requested memory, filled with copies of `PYMEM_CLEANBYTE`, used to catch reference to uninitialized memory. When a `realloc`-like function is called requesting a larger memory block, the new excess bytes are also filled with `PYMEM_CLEANBYTE`. When a `free`-like function is called, these are overwritten with `PYMEM_DEADBYTE`, to catch reference to freed memory. When a `realloc`-like function is called requesting a smaller memory block, the excess old bytes are also filled with `PYMEM_DEADBYTE`.

$p[N:N+S]$

Copies of `PYMEM_FORBIDDENBYTE`. Used to catch over- writes and reads.





Only used if the `PYMEM_DEBUG_SERIALNO` macro is defined (not defined by default).

A serial number, incremented by 1 on each call to a malloc-like or realloc-like function. Big-endian `size_t`. If “bad memory” is detected later, the serial number gives an excellent way to set a breakpoint on the next run, to capture the instant at which this block was passed out. The static function `bumpserialno()` in `obmalloc.c` is the only place the serial number is incremented, and exists so you can set such a breakpoint easily.

A realloc-like or free-like function first checks that the `PYMEM_FORBIDDENBYTE` bytes at each end are intact. If they’ve been altered, diagnostic output is written to `stderr`, and the program is aborted via `Py_FatalError()`. The other main failure mode is provoking a memory error when a program reads up one of the special bit patterns and tries to use it as an address. If you get in a debugger then and look at the object, you’re likely to see that it’s entirely filled with `PYMEM_DEADBYTE` (meaning freed memory is getting used) or `PYMEM_CLEANBYTE` (meaning uninitialized memory is getting used).

*Changed in version 3.6:* The `PyMem_SetupDebugHooks()` function now also works on Python compiled in release mode. On error, the debug hooks now use `tracemalloc` to get the traceback where a memory block was allocated. The debug hooks now also check if the GIL is held when functions of `PYMEM_DOMAIN_OBJ` and `PYMEM_DOMAIN_MEM` domains are called.

*Changed in version 3.8:* Byte patterns `0xCB` (`PYMEM_CLEANBYTE`), `0xDB` (`PYMEM_DEADBYTE`) and `0xFB` (`PYMEM_FORBIDDENBYTE`) have been replaced with `0xCD`, `0xDD` and `0xFD` to use the same values than Windows CRT debug `malloc()` and `free()`.

## The pymalloc allocator

Python has a *pymalloc* allocator optimized for small objects (smaller or equal to 512 bytes) with a short lifetime. It uses memory mappings called “arenas” with a fixed size of 256 KiB. It falls back to `PyMem_RawMalloc()` and `PyMem_RawRealloc()` for allocations larger than 512 bytes.

*pymalloc* is the **default allocator** of the `PYMEM_DOMAIN_MEM` (ex: `PyMem_Malloc()`) and `PYMEM_DOMAIN_OBJ` (ex: `PyObject_Malloc()`) domains.

The arena allocator uses the following functions:

- `VirtualAlloc()` and `VirtualFree()` on Windows,
- `mmap()` and `munmap()` if available,
- `malloc()` and `free()` otherwise.

This allocator is disabled if Python is configured with the `--without-pymalloc` option. It can also be disabled at runtime using the `PYTHONMALLOC` environment variable (ex: `PYTHONMALLOC=malloc`).

### Customize pymalloc Arena Allocator

*New in version 3.4.*

#### type **PyObjectArenaAllocator**

Structure used to describe an arena allocator. The structure has three fields:

Field	Meaning
<code>void *ctx</code>	user context passed as first argument



```
void free(void *ctx, void *ptr, size_t
size)
```

free an arena

void **PyObject\_GetArenaAllocator**(PyObjectArenaAllocator \*allocator)

Get the arena allocator.

void **PyObject\_SetArenaAllocator**(PyObjectArenaAllocator \*allocator)

Set the arena allocator.

## tracemalloc C API

*New in version 3.7.*

int **PyTraceMalloc\_Track**(unsigned int domain, uintptr\_t ptr, size\_t size)

Track an allocated memory block in the [tracemalloc](#) module.

Return 0 on success, return -1 on error (failed to allocate memory to store the trace). Return -2 if tracemalloc is disabled.

If memory block is already tracked, update the existing trace.

int **PyTraceMalloc\_Untrack**(unsigned int domain, uintptr\_t ptr)

Untrack an allocated memory block in the [tracemalloc](#) module. Do nothing if the block was not tracked.

Return -2 if tracemalloc is disabled, otherwise return 0.

## Examples

Here is the example from section [Overview](#), rewritten so that the I/O buffer is allocated from the Python heap by using the first function set:

```
PyObject *res;
char *buf = (char *) PyMem_Malloc(BUFSIZ); /* for I/O */

if (buf == NULL)
    return PyErr_NoMemory();
/* ...Do some I/O operation involving buf... */
res = PyBytes_FromString(buf);
PyMem_Free(buf); /* allocated with PyMem_Malloc */
return res;
```

The same code using the type-oriented function set:

```
PyObject *res;
char *buf = PyMem_New(char, BUFSIZ); /* for I/O */

if (buf == NULL)
    return PyErr_NoMemory();
/* ...Do some I/O operation involving buf... */
res = PyBytes_FromString(buf);
PyMem_Free(buf); /* allocated with PyMem_New */
```



Note that in the two examples above, the buffer is always manipulated via functions belonging to the same set. Indeed, it is required to use the same memory API family for a given memory block, so that the risk of mixing different allocators is reduced to a minimum. The following code sequence contains two errors, one of which is labeled as *fatal* because it mixes two different allocators operating on different heaps.

```
char *buf1 = PyMem_New(char, BUFSIZ);
char *buf2 = (char *) malloc(BUFSIZ);
char *buf3 = (char *) PyMem_Malloc(BUFSIZ);
...
PyMem_Del(buf3); /* Wrong -- should be PyMem_Free() */
free(buf2);      /* Right -- allocated via malloc() */
free(buf1);      /* Fatal -- should be PyMem_Del() */
```

In addition to the functions aimed at handling raw memory blocks from the Python heap, objects in Python are allocated and released with `PyObject_New()`, `PyObject_NewVar()` and `PyObject_Del()`.

These will be explained in the next chapter on defining and implementing new object types in C.

"